

Accelerating Entropy-Based Transformer Calibration



Michael Hu
Advisor: Karthik Narasimhan

1. Language model: $\widehat{Pr}(w_t|w_{<t})$

1. Language model: $\widehat{Pr}(w_t|w_{<t})$
2. Entropy: a measure of the amount of information being produced by a source.

1. Language model: $\widehat{Pr}(w_t|w_{<t})$
2. Entropy: a measure of the amount of information being produced by a source.
3. Entropy rate: the average amount of information per word. For a sentence, sum the entropies and divide by number of words.

1. Language model: $\widehat{Pr}(w_t|w_{<t})$
2. Entropy: a measure of the amount of information being produced by a source.
3. Entropy rate: the average amount of information per word. For a sentence, sum the entropies and divide by number of words.
 - ▶ "A loud and boisterous carnival is outside my window." \Rightarrow High entropy rate.

1. Language model: $\widehat{Pr}(w_t|w_{<t})$
2. Entropy: a measure of the amount of information being produced by a source.
3. Entropy rate: the average amount of information per word. For a sentence, sum the entropies and divide by number of words.
 - ▶ "A loud and boisterous carnival is outside my window." \Rightarrow High entropy rate.
 - ▶ "Nothing is outside my window." \Rightarrow Low entropy rate.

1. Language model: $\widehat{Pr}(w_t|w_{<t})$
2. Entropy: a measure of the amount of information being produced by a source.
3. Entropy rate: the average amount of information per word. For a sentence, sum the entropies and divide by number of words.
 - ▶ "A loud and boisterous carnival is outside my window." \Rightarrow High entropy rate.
 - ▶ "Nothing is outside my window." \Rightarrow Low entropy rate.
4. Entropy blowup: when the entropy rate of text produced by a source increases over time.

Entropy blowup does not occur in human-generated text.

- ▶ Intuition: people don't say increasingly weird things as you talk to them.

Entropy blowup does not occur in human-generated text.

- ▶ Intuition: people don't say increasingly weird things as you talk to them.

We observe entropy blowup on SOTA language models such as GPT-2.

Develop a computationally efficient technique for minimizing entropy blowup in language models.

Entropy:

$$H(Pr) := \mathbb{E}_{w_{1:t} \sim Pr} \left[\log \frac{1}{Pr(w_{1:t})} \right]$$

Entropy:

$$H(Pr) := \mathbb{E}_{w_{1:t} \sim Pr} \left[\log \frac{1}{Pr(w_{1:t})} \right]$$

One-step lookahead entropy:

$$H(\widehat{W}_{t+1} | w_{\leq t}) = \mathbb{E}_{w_{t+1} \sim \widehat{Pr}} \left[\log \frac{1}{\widehat{Pr}(w_{t+1} | w \leq t)} \right]$$

Stepwise calibration:

1. Obtain the learned language model \widehat{Pr} .

Stepwise calibration:

1. Obtain the learned language model \widehat{Pr} .
2. Scale \widehat{Pr} by α times the one-step lookahead entropy.
Renormalize to obtain a new probability distribution Pr_α .

Stepwise calibration:

1. Obtain the learned language model \widehat{Pr} .
2. Scale \widehat{Pr} by α times the one-step lookahead entropy.
Renormalize to obtain a new probability distribution Pr_α .
3. Choose α such that it minimizes the cross-entropy loss between the context and Pr_α .

Stepwise calibration:

1. Obtain the learned language model \widehat{Pr} .
2. Scale \widehat{Pr} by α times the one-step lookahead entropy.
Renormalize to obtain a new probability distribution Pr_α .
3. Choose α such that it minimizes the cross-entropy loss between the context and Pr_α .
4. Use Pr_α to predict the next word.

Computational complexity:

- ▶ V = size of vocabulary = $\sim 50,000$ in practice
- ▶ C = number of contexts = for a sentence, the number of words $\Rightarrow \sim 10$.

Computational complexity:

- ▶ V = size of vocabulary = $\sim 50,000$ in practice
- ▶ C = number of contexts = for a sentence, the number of words $\Rightarrow \sim 10$.

The one-step lookahead entropy takes $O(VC)$ time to compute. The method, as proposed, is slow in practice. The factor we can control here is V . (If you want to feed the model more context, it is your right to do so.)

The major contribution of this IW: we can approximate the true one step lookahead entropy by calculating lookahead entropies for the top 100 or so words in the vocabulary.

The major contribution of this IW: we can approximate the true one step lookahead entropy by calculating lookahead entropies for the top 100 or so words in the vocabulary.

- ▶ Top 100 words capture about 90% of the probability mass!

The major contribution of this IW: we can approximate the true one step lookahead entropy by calculating lookahead entropies for the top 100 or so words in the vocabulary.

- ▶ Top 100 words capture about 90% of the probability mass!
- ▶ $V = 50,000 \Rightarrow 100$. If the model used to take 8 hours to run, it now takes 1 minute.